

Pittsburgh, Pennsylvania412.287.8640russelldudek@gmail.comlinkedin.com/in/russelldudekgithub.com/russelldudek/Quick

MEMBER OF TECHNICAL STAFF - ENTERPRISE AI

Make the model sit on an operable backplane.

The thesis

Enterprise AI scales when outcome, context, model behavior, runtime integration, human authority, and evidence share one production contract. The weakest interface becomes the system limit.

Nominal role vs. actual mandate

The posting says Develop and optimize AI models; build ML pipelines; integrate AI into APIs and production systems; select frameworks; document architecture and deployment procedures.	The enterprise mandate likely is Turn technically capable AI into a scalable, reliable, governed operating system whose interfaces, controls, evidence, recovery, and ownership remain legible under production pressure.
---	---

Opportunity context

The staffing firm is identified, but the end client, domain, architecture, team, and current production estate are not. This brief therefore separates documented requirements from hypotheses for discovery.

Operating tensions

Model velocity vs. operability Fast experiments matter; interfaces and recovery determine trust.	Performance vs. economics Quality sits beside latency, throughput, infrastructure cost, and human effort.	Automation vs. authority Preserve judgment at exceptions, approvals, and policy boundaries.
Reuse vs. context Common patterns create leverage while workflow context remains specific.	Scale vs. recovery Reach cannot outrun monitoring, rollback, and incident ownership.	Opacity vs. action The correct first move is discovery, not fabricated certainty.

412.287.8640 russelldudek@gmail.com linkedin.com/in/russelldudek github.com/russelldudek/Quick

OPERATING MODEL

Six contracts create one production state.

OUTCOME		CONTEXT	MODEL	RUNTIME	AUTHORITY	EVIDENCE
Contract	Question			Evidence before scale		
Outcome	Which workflow or decision changes, for whom, against what baseline?			Named owner, current behavior, success measure, consequence, stop condition.		
Context	Which sources are authoritative, current, accessible, traceable, and safe?			Provenance, versions, freshness, quality, access, conflict handling.		
Model	What behavior is optimized, evaluated, bounded, or rejected?			Use-case evaluation set, thresholds, error taxonomy, fallback.		
Runtime	What must latency, throughput, availability, cost, and recovery look like?			Dependency map, load behavior, observability, degradation, rollback.		
Authority	Which actions remain advisory, reviewed, approved, reversible, or prohibited?			Decision rights, overrides, exception ownership, escalation.		
Evidence	Who owns technical health, usefulness, incidents, learning, and retirement?			Balanced scorecard, evaluation cadence, incident review, named owners.		

Decision architecture

Build / buy / adapt / reuse Select the smallest architecture that meets the outcome and reliability envelope.	Ship / bound / hold Use evidence to choose controlled proof, further hardening, or stop.	Automate / review / prohibit Make authority proportional to consequence and reversibility.
Standardize / localize Reuse proven interfaces while preserving workflow-specific judgment.	Scale / roll back / retire Expand only when the weakest interface has control and recovery.	Escalate / learn Turn incidents and user evidence into changes.

PROOF MAP

Adjacent evidence, honestly bounded.

Evidence	Transferable relevance	Claim boundary
Vape-Jet	AI-first operating transformation across robotics, machine vision, telemetry, ERP, CRM, issue flow, production, quality, support, and engineering.	Direct operating integration; not ownership of the underlying ML platform.
DudeWorth	AI readiness, structured context, retrieval, agentic workflows, human review, evaluation, governance, and reuse.	Direct design/advisory work; concepts remain concepts.
Compunetics	High-reliability engineering, manufacturing, quality, root cause, process controls, and complex technical delivery.	Advanced electronics, not modern foundation-model research.
Amazon	24/7 launch, visibility, escalation, standard work, throughput, and customer consequence at scale.	Direct operations leadership with accurately framed innovation contributions.
ZeusVu	Regulated autonomous operations and TensorFlow-based computer-vision concepts.	Concept work is not a deployed enterprise vision platform.

Strongest reasonable objection

Objection: Russell is not a conventional five-year foundation-model engineer or recent PyTorch/JAX specialist. **Response:** The fit is strongest when the client needs a hands-on systems thinker who can connect AI models to APIs, workflow, controls, reliability, evidence, ownership, and cross-functional execution.

Executive questions

1. Where is the role's center of gravity: model development, platform engineering, enterprise integration, or technical leadership?
2. Which production failure is currently most expensive?
3. What evidence distinguishes an impressive proof from a system approved to scale?
4. Which decisions belong to this role and which are shared?
5. How are model, runtime, workflow, and outcome ownership handled after deployment?

Source basis

The role description and referral URL were supplied by the candidate. Candidate claims come from the verified evidence ledger; client-specific architecture and priorities remain hypotheses for discovery.